

House Price Prediction Model - Report & Documentation

1. Objective

The goal of this project was to develop and deploy a machine learning model capable of predicting house prices based on various attributes. The process involved data preprocessing, model training and evaluation, inference testing, and deployment using FastAPI.

2. Dataset

- Source: [Kaggle House Data](#)

3. Data Preprocessing

Exploratory Data Analysis (EDA)

- **Size:** 4600 rows, 18 columns (13 numerical, 5 categorical)
- Examined correlation between features and price
- Visualized the correlation and outliers using Boxplot and Heatmaps
- **Observations:**
 - No missing values or duplicates
 - Presence of significant outliers
 - Some features moderately correlated with price

Feature Engineering

- **Location Encoding:** Categorical columns (`street` , `city` , `statezip` , `country`) were converted into geographic coordinates using the **Azure Maps API**.
- **Clustering:**
 - Optimal number of clusters determined (k=10) using silhouette score.
 - Created a new categorical feature `area_category` from cluster labels.
 - One-hot encoded `area_category` and merged into the main dataset.
- **Outlier Removal:**
 - IQR method was used to detect and remove extreme outliers.
 - Final dataset size reduced to **3727 rows**.
- **New Features:**
 - `yr_updated` : Maximum of `yr_built` and `yr_renovated`.
 - `is_renovated` : Binary indicator if house was renovated.
- **Feature Selection & Scaling:**
 - Relevant features selected for training.
 - MinMaxScaler applied for normalization.

4. Model Training & Evaluation

- **Algorithm Selection:**
 - Tried multiple regression models.
 - **GradientBoostingRegressor** provided the best results.
- **Hyperparameter Tuning:**
 - GridSearchCV with 5-fold cross-validation.
 - Found optimal parameters for best performance.
- **Train/Test Split:**
 - Data split into **85% training / 15% testing**.
 - Model trained using optimal parameters.
- **Evaluation Metrics:**
 - R² score - **75.88%**
 - Mean Absolute Error (MAE) - **72648**, which is **35.08%** of standard deviation
 - Mean Squared Error (MSE) - **10919533397**
- **Visualization:**
 - Scatterplots and error analysis plots used to visualize model performance.
- **Final Model Training:**
 - Model retrained on the full dataset.
 - Model, scalars, and coordinate mapping exported as `.pkl` files.

5. Inference Testing

- Model and pre-processing components loaded independently.
- Custom input values tested to ensure inference correctness.
- Functions modularized for clean and efficient inference.

6. Model Deployment

Framework & API

- **FastAPI** used for building the REST API.
- **Pydantic** for request validation.
- **Uvicorn** for running the application.

API & Web Interface

- **Schema:**
 - Defined `Pydantic BaseModel` schemas for input validation.
- **Prediction Flow:**
 - User inputs values via a web form.
 - JavaScript sends data as JSON to `/predict` endpoint.
 - API processes input, applies transformations, and returns a price prediction.
 - JavaScript updates the webpage with the predicted price.
- **Frontend:**
 - Clean HTML, CSS, and JavaScript template.
 - Input validation in the frontend to restrict invalid values.

Deployment

- **Local Testing:**
 - API successfully tested on localhost.
- **Hosting:**
 - Application deployed on **Render**.
 - Public URL: [House Price Prediction App](#)

Additional Notes

- Running the project **locally** requires setting up the `AZURE_MAPS_API_KEY` environment variable.
 - Hosted version on **Render** can be accessed without API key requirements.
 - **Future Improvements:**
 - Implement logging and improve error handling.
 - Integrate **DVC** or **MLflow** for model versioning.
-

Conclusion

This project successfully demonstrates the complete lifecycle of an ML model—from data preprocessing and training to deployment. The **GradientBoostingRegressor** model performed best, with an R2 score upto **75.88%** and was deployed as a web-based FastAPI application, providing users with a seamless way to predict house prices based on input attributes.